

Good strategies with n -bit memory

Nikoleta E. Glynnatsi, Christian Hilbe, Martin Nowak

We are interested in extending the results of [Akin, 2016] to strategy spaces with $n > 1$ rounds of memory. In the following, we outline our setup and our main conjecture so far.

Repeated donation game. We consider the infinitely repeated games among two players, player p and player q . Each round, they engage in the donation game with payoff matrix

$$\begin{pmatrix} b-c & -c \\ b & 0 \end{pmatrix}. \quad (1)$$

Here b and c denote the benefit and the cost of cooperation, respectively. We assume $b > c > 0$ throughout. Therefore, the payoff matrix (1) is a special case of a prisoner's dilemma.

Memory- n strategies. We assume in the following, that the players' decisions only depend on the outcome of the previous n rounds. To this end, an n -history for player p is a string $h^p = (a_{-1}^p, \dots, a_{-n}^p) \in \{C, D\}^n$. An entry a_{-k}^p corresponds to player p 's action k rounds ago. Let H^p denote the space of all n -histories of player p . Analogously, we define H^q as the set of n -histories h^q of player q . A pair $h = (h^p, h^q)$ is called an n -history of the game. We use $H = H^p \times H^q$ to denote the space of all such histories. This set contains $|H| = 2^{2n}$ elements. A *memory- n* strategy is a vector $\mathbf{p} = (p_h)_{h \in H} \in [0, 1]^{2^{2n}}$. Each entry p_h corresponds to the player's cooperation probability in the next round, depending on the outcome of the previous n rounds. One special case of such a memory- n strategy is the *round- k -repeat strategy* for some $1 \leq k \leq n$. Player p uses a *round- k -repeat strategy* $\mathbf{p}^{k\text{-Rep}}$ if in any given round, the player chooses the same action as k rounds ago. That is, if the game's n -history is such that $a_{-k}^p = C$, then $p_h^{k\text{-Rep}} = 1$; otherwise $p_h^{k\text{-Rep}} = 0$.

If the two players use memory- n strategies \mathbf{p} and \mathbf{q} , one can represent the interaction as a Markov chain with a $2^{2n} \times 2^{2n}$ transition matrix M . Let $\mathbf{v} = (v_h)_{h \in H}$ be an invariant distribution of this Markov chain. With the same method as in [Akin, 2016], one can show *Akin's Lemma*: For each k with $1 \leq k \leq n$, the invariant distribution \mathbf{v} satisfies the following relationship,

$$\mathbf{v} \cdot (\mathbf{p} - \mathbf{p}^{k\text{-Rep}}) = \sum_{h \in H} v_h (p_h - p_h^{k\text{-Rep}}) = 0. \quad (2)$$

The intuition for this result is that $\mathbf{v} \cdot \mathbf{p}$ and all $\mathbf{v} \cdot \mathbf{p}^{k\text{-Rep}}$ are just different (but equivalent) expressions for player p 's average cooperation rate. For example, $\mathbf{v} \cdot \mathbf{p}$ corresponds to a setup in which one first draws a history h according to the invariant distribution \mathbf{v} ; then one takes player p 's probability p_h to cooperate in the next round; the expectation of this procedure is $\sum_{h \in H} v_h p_h$.

Based on the invariant distribution \mathbf{v} , we can also compute the players' payoffs. To this end, let $\mathbf{S}^k = (S_h^k)_{h \in H}$ denote the vector that returns for each h the one-shot payoff that player p obtained k rounds ago,

$$S_h^k = \begin{cases} b-c & \text{if } a_{-k}^p = C \text{ and } a_{-k}^q = C \\ -c & \text{if } a_{-k}^p = C \text{ and } a_{-k}^q = D \\ b & \text{if } a_{-k}^p = D \text{ and } a_{-k}^q = C \\ 0 & \text{if } a_{-k}^p = D \text{ and } a_{-k}^q = D \end{cases} \quad (3)$$

Then we can define player p 's repeated-game payoff $s_{\mathbf{p}}$ as

$$s_{\mathbf{p}} = \mathbf{v} \cdot \mathbf{S}^1 = \mathbf{v} \cdot \mathbf{S}^2 = \dots = \mathbf{v} \cdot \mathbf{S}^n. \quad (4)$$

The equalities $\mathbf{v} \cdot \mathbf{S}^1 = \dots = \mathbf{v} \cdot \mathbf{S}^n$ correspond to the intuition that it does not matter which of the past n rounds we use to define average payoffs (this is a direct consequence of Akin's Lemma). The payoff $s_{\mathbf{q}}$ of player q can be defined analogously.

Akin's lemma imposes some natural restrictions on which invariant distributions \mathbf{v} are possible. Similar to the paper by [Akin, 2016], we hope to exploit these restrictions to characterise the good memory- n strategies (to be defined below).

Good strategies. We say $h = (h^p, h^q)$ is the mutual cooperation history if $h^p = h^q = (C, \dots, C)$. A memory- n strategy \mathbf{p} is called agreeable if it prescribes to cooperate with probability 1 after the mutual cooperation history. The strategy \mathbf{p} is called good if it is agreeable and if expected payoffs satisfy

$$s_{\mathbf{q}} \geq b - c \quad \Rightarrow \quad s_{\mathbf{q}} = s_{\mathbf{p}} = b - c. \quad (5)$$

We wish to characterise all good memory- n strategies of the repeated donation game. To start with, in the following we begin with the simplest non-trivial case.

The case of 2-bit reactive strategies. We say a memory- n strategy \mathbf{p} is n -bit reactive if it only depends on the co-player's n -history (independent of the focal player's own actions during the past n rounds). Formally, $\mathbf{p} = (p_h)_{h \in H}$ is n -bit reactive if for any two histories $h = (h^p, h^q)$ and $\tilde{h} = (\tilde{h}^p, \tilde{h}^q)$ with $h^q = \tilde{h}^q$ it follows that $p_h = p_{\tilde{h}}$. In particular, for $n=2$ such a player uses at most 4 different cooperation probabilities, depending on the co-player's actions during the last 2 rounds (CC, CD, DC, DD , where the first letter refers to the co-player's action in the second-to-last round, and the second letter refers to the last round). Slightly abusing notation, we write 2-bit reactive strategies as

$$\hat{\mathbf{p}} = (\hat{p}_{CC}, \hat{p}_{CD}, \hat{p}_{DC}, \hat{p}_{DD}). \quad (6)$$

For 2-bit reactive strategies, we have the following conjecture.

Conjecture. Let $\hat{\mathbf{p}}$ be an agreeable 2-bit reactive strategy, i.e. $\hat{p}_{CC} = 1$. The following are equivalent

(i) The strategy $\hat{\mathbf{p}}$ is good.

(ii) The entries of $\hat{\mathbf{p}}$ satisfy $\hat{p}_{DD} < 1 - \frac{c}{b}$ and $\frac{\hat{p}_{CD} + \hat{p}_{DC}}{2} < 1 - \frac{c}{2b}$.

Our evidence for this conjecture is as follows. The direction (i) \Rightarrow (ii) is straightforward. If \mathbf{p} is good, it needs to satisfy condition (5) for all \mathbf{q} . In particular, the condition needs to be satisfied when \mathbf{q} is either *ALLD* (the strategy that always defects), or *Alternator* (the strategy that cooperates if and only if it didn't cooperate the previous round). By checking these two strategies explicitly, one gets (ii).

The direction (ii) \Rightarrow (i) we could not prove yet. However, we have strong numerical evidence (see Figure next page). We have sampled 10^4 random agreeable 2-bit reactive strategies and checked numerically whether or not they are Nash equilibria. We found that exactly those strategies are Nash equilibria that satisfy the conditions in (ii).

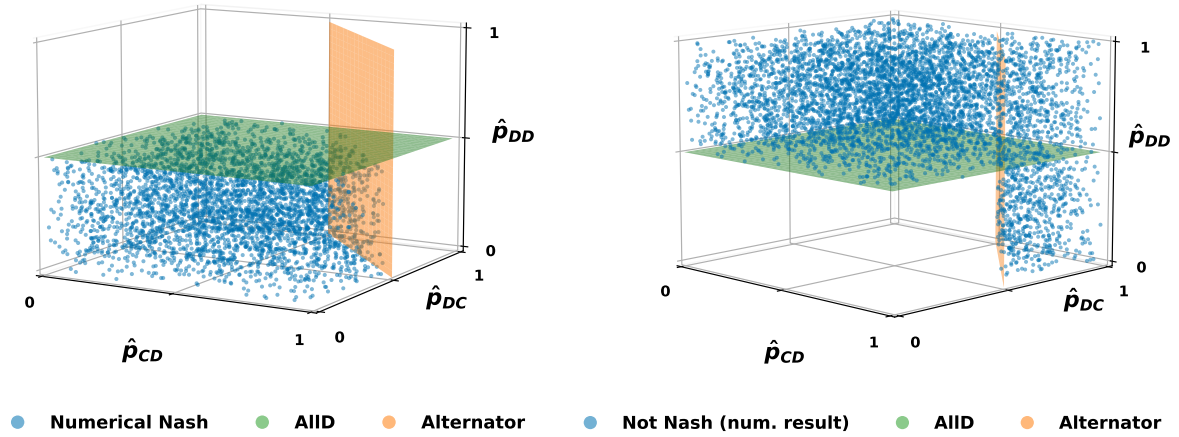


Figure 1: We generated 10^4 agreeable 2-bit strategies of the form $\hat{\mathbf{p}} = (1, \hat{p}_{CD}, \hat{p}_{DC}, \hat{p}_{DD})$ uniformly at random. For each such $\hat{\mathbf{p}}$ we numerically checked whether or not the strategy is a Nash equilibrium. To this end, by an argument similar to the one in [Press and Dyson, 2012] and [McAvoy and Nowak, 2019], it is sufficient to check all deviations towards pure memory-2 strategies \mathbf{q} . If a strategy $\hat{\mathbf{p}}$ was numerically found to be a Nash equilibrium, we depict it as a blue dot in the left panel. Otherwise, if it was found not to be a Nash equilibrium, we depict it as a blue dot in the right panel. Points below the green plane satisfy $\hat{p}_{DD} < 1 - \frac{\epsilon}{b}$. Points left to the orange plane satisfy $\hat{p}_{CD} + \hat{p}_{DC} < 2 - \frac{\epsilon}{b}$. We find that all Nash equilibria satisfy these two inequalities (left panel). Conversely, our numerical results suggest that all $\hat{\mathbf{p}}$ that are not Nash equilibria violate at least one of these inequalities (right panel). Parameters: $b=2$, $c=1$.

References

- E. Akin. The iterated prisoner’s dilemma: good strategies and their dynamics. *Ergodic Theory, Advances in Dynamical Systems*, pages 77–107, 2016.
- A. McAvoy and M. A. Nowak. Reactive learning strategies for iterated games. *Proceedings of the Royal Society A*, 475(2223):20180819, 2019.
- W. H. Press and F. J. Dyson. Iterated prisoner’s dilemma contains strategies that dominate any evolutionary opponent. *Proceedings of the National Academy of Sciences*, 109(26):10409–10413, 2012.