Good strategies with n-bit memory

Nikoleta E. Glynatsi, Christian Hilbe, Martin Nowak

We are interested in extending the results of [Akin, 2016] to strategy spaces with $n > 1$ rounds of memory. In the following, we outline our setup and our main conjecture so far.

Repeated donation game. We consider the infinitely repeated games among two players, player p and player q. Each round, they engage in the donation game with payoff matrix

$$\begin{pmatrix} b-c & -c \\ b & 0 \end{pmatrix}. \tag{1}$$

Here b and c denote the benefit and the cost of cooperation, respectively. We assume $b > c > 0$ throughout. Therefore, the payoff matrix (1) is a special case of a prisoner's dilemma.

Memory-n strategies. We assume in the following, that the players' decisions only depend on the outcome of the previous n rounds. To this end, an *n-history for player p* is a string $h^p = (a^p_{-1}, \dots, a^p_{-n}) \in \{C, D\}^n$. An entry a^p_{-k} corresponds to player p's action k rounds ago. Let H^p denote the space of all n-histories of player p. Analogously, we define H^q as the set of n-histories h^q of player q. A pair $h = (h^p, h^q)$ is called an *n-history of the game*. We use $H = H^p \times H^q$ to denote the space of all such histories. This set contains $|H| = 2^{2n}$ elements. A *memory-n* strategy is a vector $\mathbf{p} = (p_h)_{h\in H} \in [0,1]^{2n}$. Each entry p_h corresponds to the player's cooperation probability in the next round, depending on the outcome of the previous n rounds. One special case of such a memory-n strategy is the *round-k-repeat strategy*. Player p uses a *round-k-repeat strategy* $\mathbf{p}^{k-\mathrm{Rep}}$ if in any given round, the player chooses the same action as k rounds ago. That is, if the game's n-history is such that $a^p_{-k} = C$, then $p_h^{k-\mathrm{Rep}} = 1$; otherwise $p_h^{k-\mathrm{Rep}} = 0$.

If the two players use memory-n strategies $\mathbf{p}$ and $\mathbf{q}$, one can represent the interaction as a Markov chain with a $2^{2n} \times 2^{2n}$ transition matrix M. Let $\mathbf{v} = (v_h)_{h\in H}$ be an invariant distribution of this Markov chain. With the same method as in [Akin, 2016], one can show *Akin's Lemma*: For each k with $1 \le k \le n$, the invariant distribution $\mathbf{v}$ satisfies the following relationship,

$$\mathbf{v} \cdot (\mathbf{p} - \mathbf{p}^{k-\mathrm{Rep}}) = \sum_{h\in H} v_h (p_h - p_h^{k-\mathrm{Rep}}) = 0. \tag{2}$$

The intuition for this result is that $\mathbf{v}\cdot\mathbf{p}$ and all $\mathbf{v}\cdot\mathbf{p}^{k-\mathrm{Rep}}$ are just different (but equivalent) expressions for player p's average cooperation rate. For example, $\mathbf{v}\cdot\mathbf{p}$ corresponds to a setup in which one first draws a history h according to the invariant distribution $\mathbf{v}$; then one takes player p's probability p_h to cooperate in the next round; the expectation of this procedure is $\sum_{h\in H} v_h p_h$.

Based on the invariant distribution $\mathbf{v}$, we can also compute the players' payoffs. To this end, let $\mathbf{S}^k = (S^k_h)_{h\in H}$ denote the vector that returns for each h the one-shot payoff that player p obtained k rounds ago,

$$S^k_h = \begin{cases} b-c & \text{if } a^p_{-k} = C \text{ and } a^q_{-k} = C \\ -c & \text{if } a^p_{-k} = C \text{ and } a^q_{-k} = D \\ b & \text{if } a^p_{-k} = D \text{ and } a^q_{-k} = C \\ 0 & \text{if } a^p_{-k} = D \text{ and } a^q_{-k} = D \end{cases} \tag{3}$$

Then we can define player p's repeated-game payoff $s_\mathbf{p}$ as

$$s_\mathbf{p} = \mathbf{v} \cdot \mathbf{S}^1 = \mathbf{v} \cdot \mathbf{S}^2 = \ldots = \mathbf{v} \cdot \mathbf{S}^n. \tag{4}$$

The equalities $\mathbf{v} \cdot \mathbf{S}^1 = \ldots = \mathbf{v} \cdot \mathbf{S}^n$ correspond to the intuition that it does not matter which of the past n rounds we use to define average payoffs (this is a direct consequence of Akin's Lemma). The payoff $s_\mathbf{q}$ of player q can be defined analogously.

Akin's lemma imposes some natural restrictions on which invariant distributions $\mathbf{v}$ are possible. Similar to the paper by [Akin, 2016], we hope to exploit these restrictions to characterise the good memory-n strategies (to be defined below).

Good strategies. We say $h = (h^p, h^q)$ is the mutual cooperation history if $h^p = h^q = (C, \ldots, C)$. A memory-$n$ strategy $\mathbf{p}$ is called agreeable if it prescribes to cooperate with probability 1 after the mutual cooperation history. The strategy $\mathbf{p}$ is called good if it is agreeable and if expected payoffs satisfy

$$s_\mathbf{q} \geq b - c \qquad \Rightarrow \qquad s_\mathbf{q} = s_\mathbf{p} = b - c, \tag{5}$$

We wish to characterise all good memory-n strategies of the repeated donation game. To start with, in the following we begin with the simplest non-trivial case.

The case of 2-bit reactive strategies. We say a memory-n strategy $\mathbf{p}$ is *n-bit reactive* if it only depends on the co-player's n-history (independent of the focal player's own actions during the past n rounds). Formally, $\mathbf{p} = (p_h)_{h \in H}$ is n-bit reactive if for any two histories $h = (h^p, h^q)$ and $\tilde{h} = (\tilde{h}^p, \tilde{h}^q)$ with $h^q = \tilde{h}^q$ it follows that $p_h = p_{\tilde{h}}$. In particular, for $n = 2$ such a player uses at most 4 different cooperation probabilities, depending on the co-player's actions during the last 2 rounds (CC, CD, DC, DD, where the first letter refers to the last round, and the second letter refers to the second-to-last round). Slightly abusing notation, we write 2-bit reactive strategies as

$$\hat{\mathbf{p}} = (\hat{p}_{CC}, \hat{p}_{CD}, \hat{p}_{DC}, \hat{p}_{DD}). \tag{6}$$

For 2-bit reactive strategies, we have the following conjecture.

Conjecture. Let $\hat{\mathbf{p}}$ be an agreeable 2-bit reactive strategy, i.e. $\hat{p}_{CC} = 1$. The following are equivalent

(i) The strategy $\hat{\mathbf{p}}$ is good.

(ii) The entries of $\hat{\mathbf{p}}$ satisfy $\hat{p}_{DD} < 1 - \dfrac{c}{b}$ and $\dfrac{\hat{p}_{CD} + \hat{p}_{DC}}{2} < 1 - \dfrac{c}{2b}$.

Our evidence for this conjecture is as follows. The direction (i) $\Rightarrow$ (ii) is straightforward. If $\mathbf{p}$ is good, it needs to satisfy condition (5) for all $\mathbf{q}$. In particular, the condition needs to be satisfied when $\mathbf{q}$ is either *ALLD* (the strategy that always defects), or *Alternator* (the strategy that cooperates if and only if it didn't cooperate the previous round). By checking these two strategies explicitly, one gets (ii).

The direction (ii) $\Rightarrow$ (i) we could not prove yet. However, we have strong numerical evidence (see Figure next page). We have sampled 10^4 random agreeable 2-bit reactive strategies and checked numerically whether or not they are Nash equilibria. We found that exactly those strategies are Nash equilibria that satisfy the conditions in (ii).

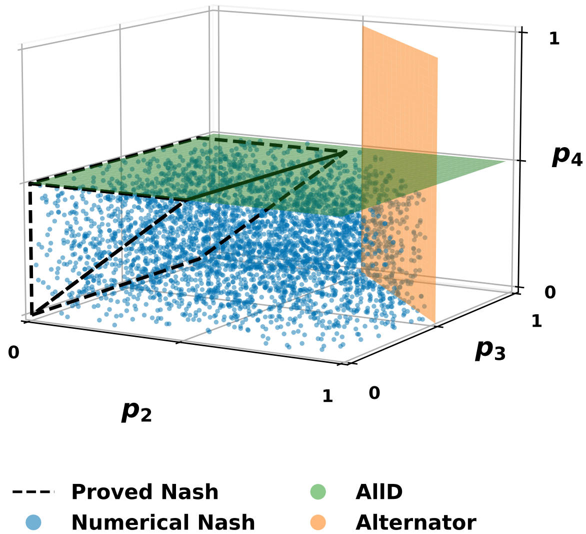

Figure 1: We generated 10^4 agreeable 2-bit strategies of the form $\hat{\mathbf{p}} = (1, \hat{p}_2, \hat{p}_3, \hat{p}_4)$ uniformly at random. For each such $\hat{\mathbf{p}}$ we numerically checked whether or not the strategy is a Nash equilibrium. To this end, by an argument similar to the one in [Press and Dyson, 2012] and [McAvoy and Nowak, 2019], it is sufficient to check all deviations towards pure memory-2 strategies $\mathbf{q}$. If a strategy $\hat{\mathbf{p}}$ was numerically found to be a Nash equilibrium, we depict it as a blue dot; otherwise we do not depict it. Points below the green plane satisfy $\hat{p}_4 \leq 1 - \frac{c}{b}$. Points left to the orange plane satisfy $\hat{p}_2 + \hat{p}_3 < 2 - \frac{c}{b}$. We find that all Nash equilibria satisfy these two conditions. Conversely, our numerical results suggest that all $\hat{\mathbf{p}}$ that satisfy these two conditions are Nash equilibria. [The black dashed lines correspond to some analytical result that we would be happy to discuss in person.] Parameters: $b=2$, $c=1$.

References

E. Akin. The iterated prisoner's dilemma: good strategies and their dynamics. *Ergodic Theory, Advances in Dynamical Systems*, pages 77–107, 2016.

A. McAvoy and M. A. Nowak. Reactive learning strategies for iterated games. *Proceedings of the Royal Society A*, 475(2223):20180819, 2019.

W. H. Press and F. J. Dyson. Iterated prisoner's dilemma contains strategies that dominate any evolutionary opponent. *Proceedings of the National Academy of Sciences*, 109(26):10409–10413, 2012.